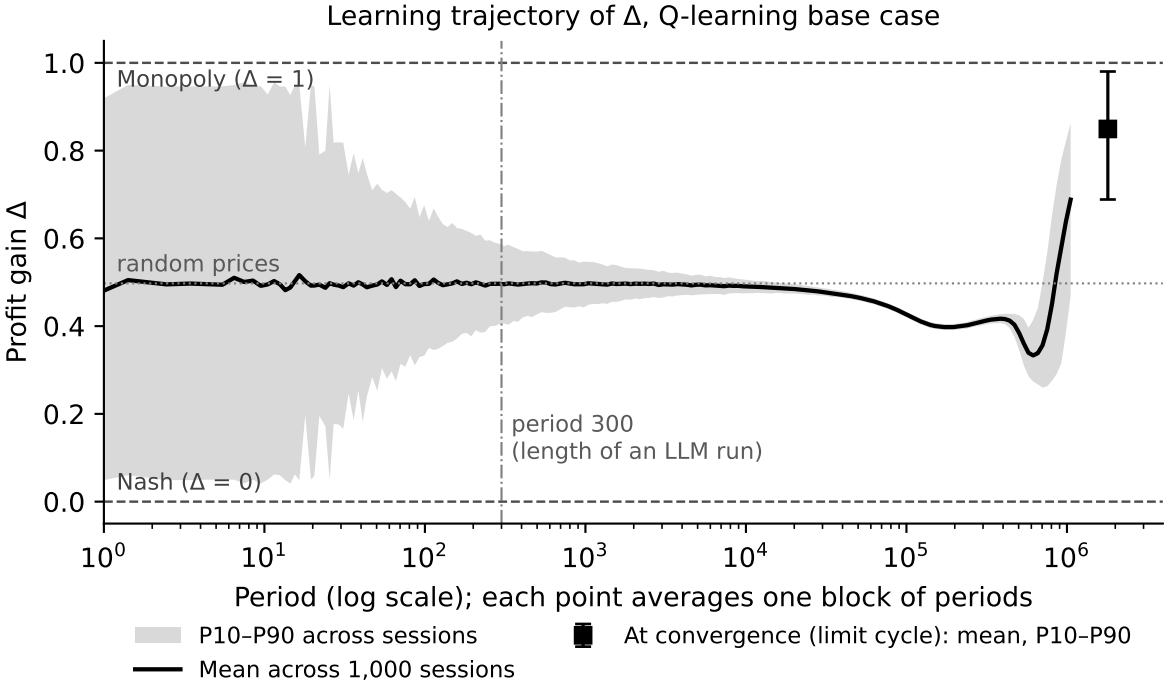
Learning trajectory of Δ, Q-learning base case
Monopoly (Δ = 1)
random prices
Nash (Δ = 0)
period 300
(length of an LLM run)
Profit gain Δ
1.0
0.8
0.6
0.4
0.2
0.0
10^0
10^1
10^2
10^3
10^4
10^5
10^6
Period (log scale); each point averages one block of periods
P10–P90 across sessions
Mean across 1,000 sessions
At convergence (limit cycle): mean, P10–P90